

Compulsory exercise 2 TMA4268: Group 6

TMA4268 Statistical Learning V2022

August Arnstad, Markus Stokkenes and Ulrik Unneberg

04 april, 2022

We would like to start by saying that in order to avoid getting a massive size of this report, some output has been commented out. We did not see it necessary to include what has been left out. We apologize for having a paper more than 14 pages long, but most of it is plots and output.

Task 1)

a)

```
#str(Boston)

set.seed(1)

# pre-processing by scaling NB! Strictly speaking, pre-processing should be done
# on a training set only and it should be done on a test set with statistics of
# the pre-processing from the training set. But, we're preprocessing the entire
# dataset here for convenience.
boston <- scale(Boston, center = T, scale = T)

#PROBLEM 1a)
#Perform Forward Stepwise Selection and Backward Stepwise Selection on boston.train method,
# and plot a graph of adjusted R2 on the y-axis and a number of predictors on the x-axis.

# split into training and test sets
train.ind = sample(1:nrow(boston), 0.8 * nrow(boston)) #We do a 80-20 split for training and test respo
boston.train = data.frame(boston[train.ind, ])
summary(boston.train )
```

##	crim	zn	indus	chas
##	Min. : -0.41937	Min. : -0.487240	Min. : -1.55630	Min. : -0.27233
##	1st Qu.: -0.41055	1st Qu.: -0.487240	1st Qu.: -0.86902	1st Qu.: -0.27233
##	Median : -0.39028	Median : -0.487240	Median : -0.21089	Median : -0.27233
##	Mean : -0.01658	Mean : 0.000275	Mean : 0.01823	Mean : 0.02003
##	3rd Qu.: 0.01197	3rd Qu.: 0.102320	3rd Qu.: 1.01499	3rd Qu.: -0.27233
##	Max. : 8.12884	Max. : 3.800473	Max. : 2.42017	Max. : 3.66477
##	nox	rm	age	dis
##	Min. : -1.46443	Min. : -3.87641	Min. : -2.333128	Min. : -1.26582
##	1st Qu.: -0.92076	1st Qu.: -0.56344	1st Qu.: -0.886356	1st Qu.: -0.80613

```
## Median :-0.14407 Median :-0.09555 Median : 0.363251 Median :-0.29410
## Mean :-0.01425 Mean : 0.02572 Mean :-0.001632 Mean :-0.00502
## 3rd Qu.: 0.59809 3rd Qu.: 0.49439 3rd Qu.: 0.906790 3rd Qu.: 0.63942
## Max. : 2.72965 Max. : 3.47325 Max. : 1.116390 Max. : 3.95660
## rad tax ptratio black
## Min. :-0.98187 Min. :-1.31269 Min. :-2.7047 Min. :-3.87923
## 1st Qu.: -0.63733 1st Qu.: -0.75495 1st Qu.: -0.4876 1st Qu.: 0.19238
## Median :-0.52248 Median :-0.43455 Median : 0.2977 Median : 0.37862
## Mean : 0.00854 Mean : 0.02023 Mean : 0.0234 Mean :-0.03297
## 3rd Qu.: 1.65960 3rd Qu.: 1.52941 3rd Qu.: 0.8058 3rd Qu.: 0.42720
## Max. : 1.65960 Max. : 1.79642 Max. : 1.6372 Max. : 0.44062
## lstat medv
## Min. :-1.52961 Min. :-1.90634
## 1st Qu.: -0.81403 1st Qu.: -0.62605
## Median :-0.18107 Median :-0.11230
## Mean :-0.01568 Mean : 0.03562
## 3rd Qu.: 0.56986 3rd Qu.: 0.32534
## Max. : 3.54526 Max. : 2.98650
```

```
boston.test = data.frame(boston[-train.ind, ])
```

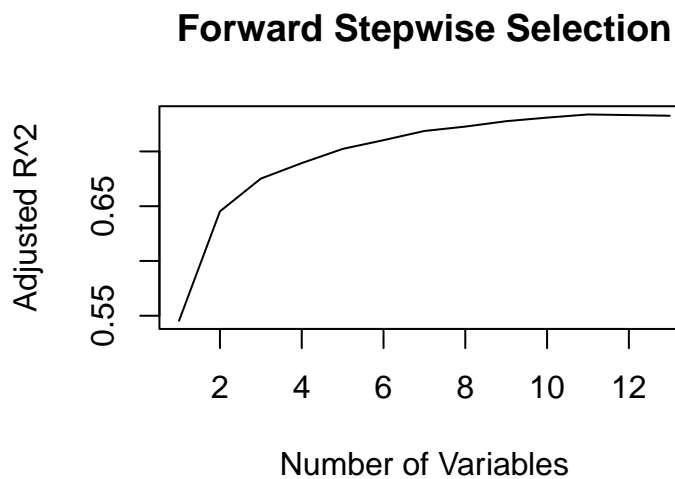
```
# Forward
```

```
regfit_fwd = regsubsets(medv~., data = boston.train, nvmax=14, method = "forward")
```

```
regfit_fwd_summary = summary(regfit_fwd)
```

```
#regfit_fwd_summary
```

```
plot(regfit_fwd_summary$adjr2, xlab = "Number of Variables", ylab = "Adjusted R^2", type = "l", main="F
```



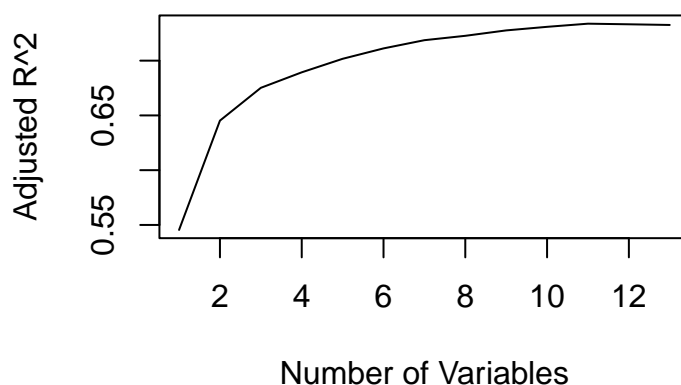
```
# Backward
```

```
regfit_bwd = regsubsets(medv~., data = boston.train, nvmax=14, method = "backward")
```

```
regfit_bwd_summary = summary(regfit_bwd)
```

```
plot(regfit_bwd_summary$adjr2, xlab = "Number of Variables", ylab = "Adjusted R^2", type = "l", main="B
```

Backward Stepwise Selection



b)

```
regfit_fwd_summary
```

```
## Subset selection object
## Call: regsubsets.formula(medv ~ ., data = boston.train, nvmax = 14,
##   method = "forward")
## 13 Variables (and intercept)
##      Forced in Forced out
## crim      FALSE      FALSE
## zn         FALSE      FALSE
## indus      FALSE      FALSE
## chas       FALSE      FALSE
## nox        FALSE      FALSE
## rm         FALSE      FALSE
## age        FALSE      FALSE
## dis        FALSE      FALSE
## rad        FALSE      FALSE
## tax        FALSE      FALSE
## ptratio    FALSE      FALSE
## black      FALSE      FALSE
## lstat      FALSE      FALSE
## 1 subsets of each size up to 13
## Selection Algorithm: forward
##      crim zn  indus chas nox rm  age dis rad tax ptratio black lstat
## 1 ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 2 ( 1 ) " " " " " " " " " " "*" " " " " " " " " " " " " " " " "
## 3 ( 1 ) " " " " " " " " " " "*" " " " " " " " " " " " " " " " "
## 4 ( 1 ) " " " " " " " " " " "*" " " "*" " " " " " " " " " " " "
## 5 ( 1 ) " " " " " " " " " " "*" " " "*" " " " " " " " " " " " "
## 6 ( 1 ) " " " " " " " " " " "*" "*" " " " "*" " " " " " " " " " "
## 7 ( 1 ) " " " " " " " "*" "*" "*" " " " "*" " " " " " " " " " "
## 8 ( 1 ) " " " " " " " "*" "*" "*" " " " "*" "*" " " " " " " " " "
```

```
## 9  ( 1 ) " " " " " " "*" "*" "*" " " " "*" "*" "*" "*" "*" "*"
## 10 ( 1 ) " " "*" " " " "*" "*" "*" " " " "*" "*" "*" "*" "*" "*"
## 11 ( 1 ) "*" "*" " " " "*" "*" "*" " " " "*" "*" "*" "*" "*" "*"
## 12 ( 1 ) "*" "*" "*" "*" "*" "*" "*" " " " "*" "*" "*" "*" "*" "*"
## 13 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" *
```

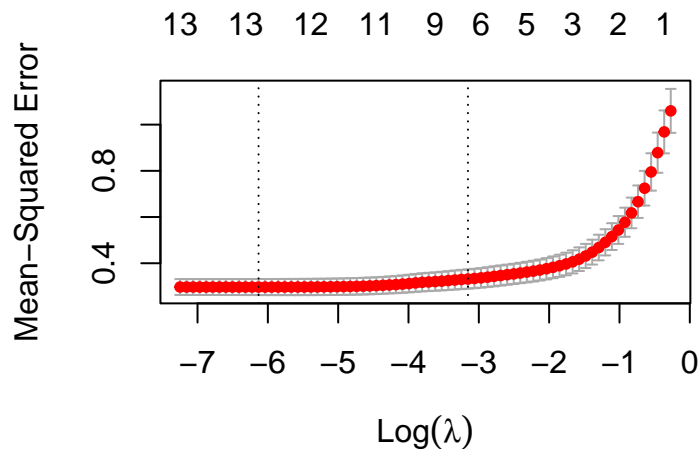
From the fourth line of the summary: The four best predictors, according to the Forward Stepwise Selection, is the predictors *rm*, *dis*, *ptratio*, and *lstat*.

c)

i) We will now do a 5-fold cross-validation on the `boston.train` data set, with the Lasso, and plot the MSE as a function of $\log \lambda$

```
set.seed(1)
x = model.matrix(medv~., data=boston.train)
y = boston.train$medv

set.seed(1)
cv.lasso = cv.glmnet(x, y , nfolds=5, alpha=1)
plot(cv.lasso)
```



ii) We λ with minimum MSE is marked with the dotted line to the left:

```
cv.lasso$lambda.min
```

```
## [1] 0.002172032
```

iii) We will now find the coefficients at the best λ , i.e. at one standard deviation away from the lowest, in the direction of fewer predictors, i.e. to the right. This is a model consisting of 7 predictors, and we find the coefficients with

```
coef(cv.lasso)
```

```
## 15 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept)  0.024709949
## (Intercept)  .
## crim        .
## zn          .
## indus       .
## chas        0.070509379
## nox         -0.004688111
## rm          0.350606336
## age         .
## dis         -0.060636734
## rad         .
## tax         .
## ptratio     -0.165115064
## black       0.081161951
## lstat       -0.423934302
```

d)

TRUE, FALSE, FALSE, TRUE

Task 2

```
library(MASS)
set.seed(1)

# load a synthetic dataset
id <- "1CWZYfrL0rFdrIZ6Hv73e3xxt0SFgU4Ph" # google file ID
synthetic <- read.csv(sprintf("https://docs.google.com/uc?id=%s&export=download",
                              id))

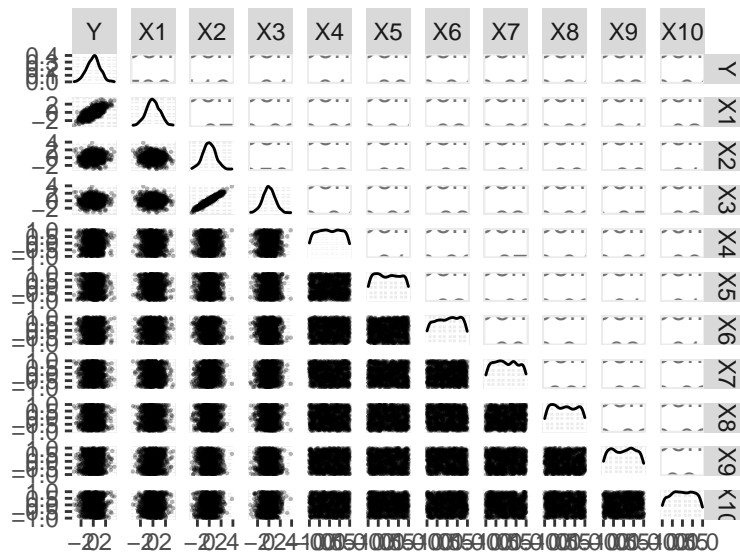
# split into training and test sets
train.ind = sample(1:nrow(synthetic), 0.8 * nrow(synthetic))
synthetic.train = data.frame(synthetic[train.ind, ])
synthetic.test = data.frame(synthetic[-train.ind, ])

# show head(...) Y: response variable; X: predictor variable
head(synthetic)
```

```
##           Y           X1           X2           X3           X4           X5
## 1 -1.43753239 -0.75905055 -0.69720326 -0.3016852 -0.7434697  0.8807558
## 2 -1.70972989 -0.28635632  0.04809182  0.5791725 -0.7446170  0.9935311
## 3  1.33931240  0.09574117 -0.89605758 -0.9636347  0.5554647 -0.5341800
## 4  0.20354906 -0.28702695  1.72952687  1.4289705 -0.1596993 -0.7161976
## 5 -0.09261896  0.02345825  0.51201583  0.1544345  0.4318039 -0.8674060
## 6  1.69952325  1.19231791 -0.98179754 -0.9567773 -0.6933918  0.4656891
```

```
##           X6           X7           X8           X9           X10
## 1 -0.8705750 -0.7448252 -0.4639697  0.62502272 -0.8149674
## 2  0.3532248 -0.5860332 -0.7964403  0.84868110 -0.1065119
## 3  0.4707434 -0.6588069 -0.7327518 -0.29429307  0.6588927
## 4 -0.7774007  0.2502145  0.5987052 -0.04428773  0.6247479
## 5 -0.9066908  0.8946086 -0.9700185  0.09082626  0.6102134
## 6 -0.7381794  0.8650175  0.4108119  0.75677429 -0.2281439
```

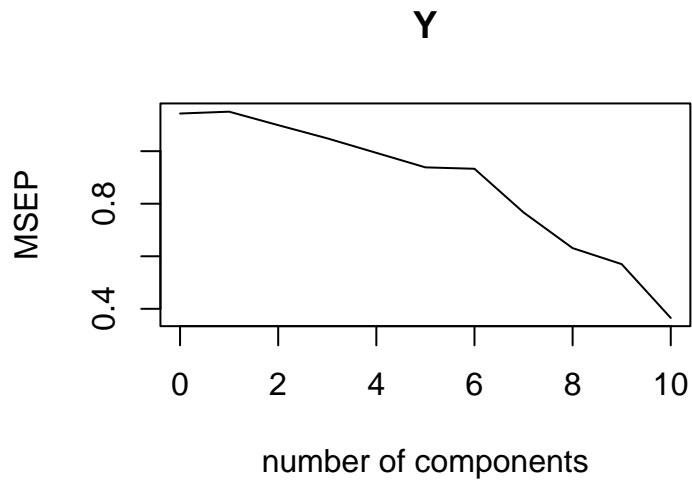
```
ggpairs(synthetic, lower = list(continuous = wrap("points", alpha = 0.3, size=0.1),
                               combo = wrap("dot", alpha = 0.4, size=0.2) ),)
```



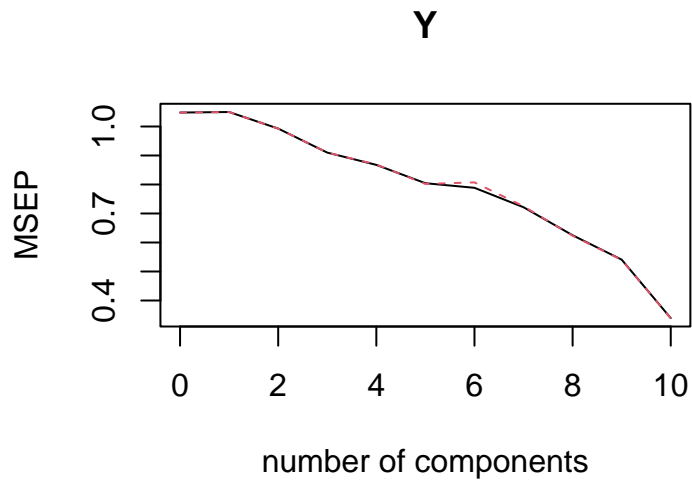
We first take a look at the data. At first glance, we first observe that there is a prominent correlation between X_2 and X_3 , and between Y and X_1 .

a)

```
#Principle Component Regression
library(pls)
set.seed(1)
pcr.fit = pcr(Y ~., data=synthetic.train, scale = TRUE, validation = "CV")
#summary(pcr.fit)
validationplot(pcr.fit, val.type = "MSEP", newdata =synthetic.test)
```

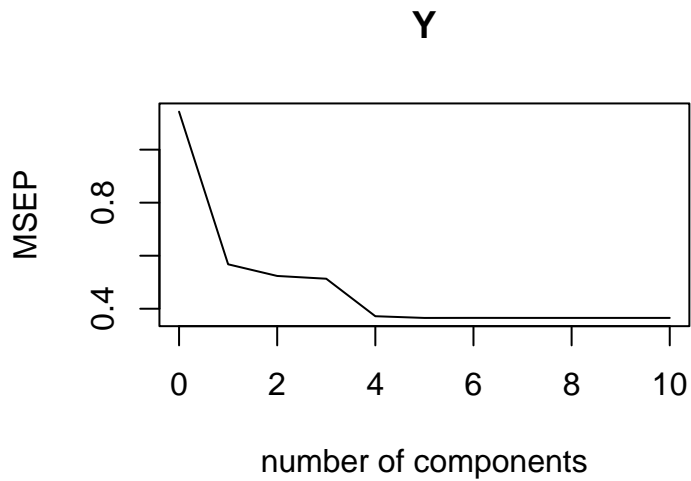


```
#Testing both cross-validation and with test data. Similar results.
validationplot(pcr.fit, val.type = "MSEP")
```

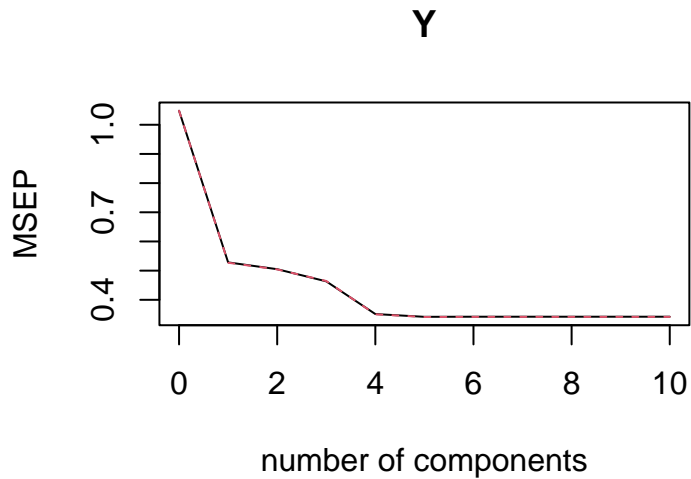


The MSEP for the PCR is clearly minimized with 10 components, i.e. a full model. This suggests that a PCR model won't benefit over an Ordinary Least Squares model.

```
#Partial Least Squares
pls.fit = plsr(Y ~., data=synthetic.train, scale = TRUE, validation = "CV")
#summary(pls.fit)
validationplot(pls.fit, val.type = "MSEP", newdata = synthetic.test)
```



```
#Testing both cross-validation and with test data. Similar results.
validationplot(pls.fit, val.type = "MSEP")
```



We observe that for the Partial Least Squares model, the MSEP flat out at the minimum MSEP from 4 components or more. This is a substantial improvement in the number of predictors; 4 instead of 10.

b)

In general, the PCR is an unsupervised method, meaning that the correlation between predictors and the response won't influence the choice of principle components. The PCR selects a linear combination of predictors along the direction with the most variance. The PLS, on the other hand, is a supervised method, prioritizing the predictors which correlates the most to Y . The components in the PLS is a linear combination of the predictors, where each of them is weighted according to how much Y varies with the particular predictor.

As we see from the ggpairs plot, the strong correlation between Y and X_1 suggests that a simple model describes the trends quite well. The PCR doesn't pick up this relation, as the model is unsupervised; the principle component might be in any direction where the variance of the predictors is large. We observe that the variables X_2 and X_3 have a large correlation, and thus common covariance along the same direction. These variables will dominate in the first principle component, completely ignoring the relationship between Y and X_1 . As there seems to be little correlation between Y and X_2, X_3 . This explains why the MSEP doesn't seem to change going from zero to one component in the PCR.

The first component in the PLS is dominated by X_1 , and we thus see that when we include this component, the MSEP drop drastically. It further drops to the minimum MSEP when the number of components reach four.

Task 3)

a)

TRUE, FALSE, FALSE, TRUE

b)

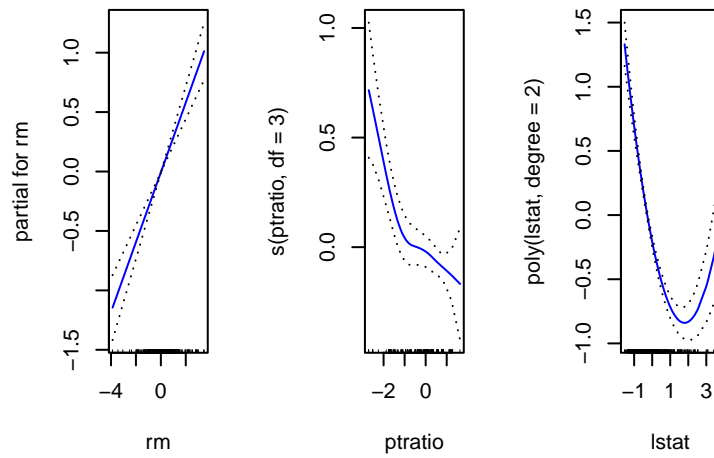
```
set.seed(1)

data(Boston)

boston<-scale(Boston, center=T, scale=T)

train.id=sample(1:nrow(boston), 0.8*nrow(boston))
boston.train = data.frame(boston[train.id, ])
#str(boston.train)
genaddmod<-gam(medv~rm + s(ptratio, df=3) + poly(lstat, degree=2), data=boston.train)

par(mfrow=c(1,3))
plot(genaddmod, se=T, col="blue")
```

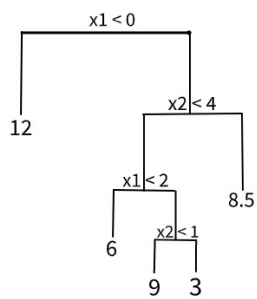


Task 4

a)

FALSE, TRUE, TRUE, TRUE

b)



c)

```
library(tidyverse)
library(palmerpenguins) # Contains the data set 'penguins'.
library(tree)
```

```

library(randomForest)
library(e1071)
data(penguins)

names(penguins) <- c("species", "island", "billL", "billD", "flipperL", "mass", "sex", "year")

Penguins_reduced <- penguins %>% dplyr::mutate(mass = as.numeric(mass), flipperL = as.numeric(flipperL))

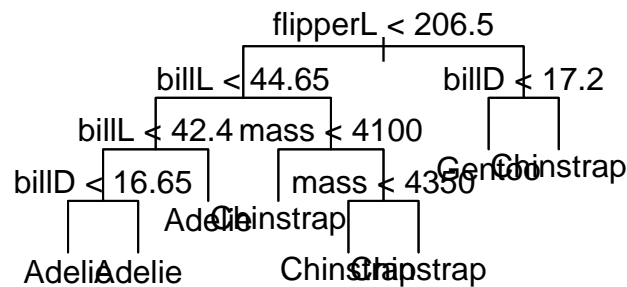
# We do not want 'year' in the data (this will not help for future predictions)
Penguins_reduced <- Penguins_reduced[, -c(8)]

set.seed(4268)
# 70% of the sample size for training set
training_set_size <- floor(0.7 * nrow(Penguins_reduced))
train_ind <- sample(seq_len(nrow(Penguins_reduced)), size = training_set_size)
train <- Penguins_reduced[train_ind, ]
test <- Penguins_reduced[-train_ind, ]

# create tree with default parameters using gini index
penguin.tree <- tree(species ~ ., train, split = "gini")

# plot the full tree
plot(penguin.tree, type = "uniform")
text(penguin.tree, pretty = 0)

```



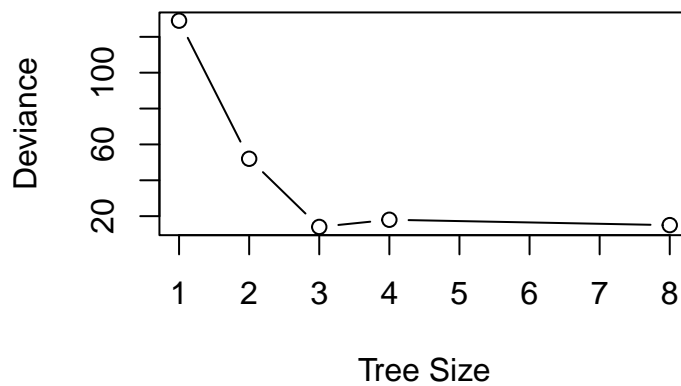
```

set.seed(123)

# cost-complexity pruning using 10-fold CV
cv.penguin <- cv.tree(penguin.tree, FUN = prune.misclass, K = 10)

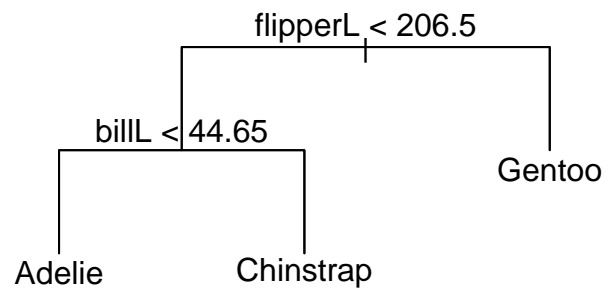
# plot deviance as a function of tree size
plot(cv.penguin$dev ~ cv.penguin$size, type = "b", xlab = "Tree Size", ylab = "Deviance")

```



```
# prune the tree according to the optimal tree size
prune.penguin <- prune.tree(penguin.tree, best = 3)

# plot the pruned tree
plot(prune.penguin, type = "uniform")
text(prune.penguin, pretty = 0)
```



```
# get predictions from training data using the pruned tree
pred.cv <- predict(prune.penguin, newdata = test, type = "class")

# create confusion table for pruned tree
misclass.cv <- table(pred.cv, test$species)

# misclassification rate
error.cv <- 1 - sum(diag(misclass.cv))/sum(misclass.cv)
```

```
cat("Misclassification error rate after cost-complexity pruning:", error.cv)
```

```
## Misclassification error rate after cost-complexity pruning: 0.06
```

d)

```
# create tree using random forest method
rf.penguin = randomForest(species ~ ., data = train, mtry = round(sqrt(ncol(Penguins_reduced) - 1)), ntrees = 1000)

# get predictions
pred.rf = predict(rf.penguin, newdata = test, type = "class")

# confusion table
misclass.rf <- table(pred.rf, test$species)

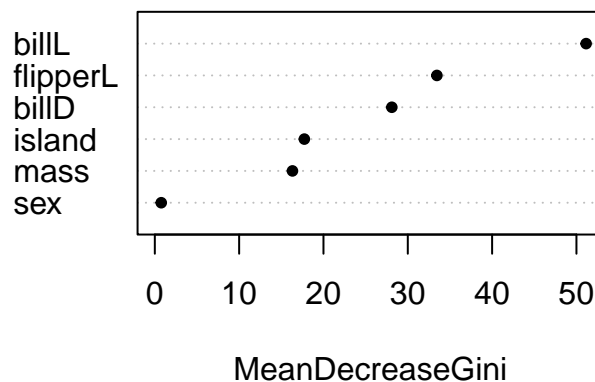
# misclassification rate
error.rf = 1 - sum(diag(misclass.rf))/sum(misclass.rf)

cat("Misclassification error rate for random forest method:", error.rf)
```

```
## Misclassification error rate for random forest method: 0.02
```

Since we are doing classification we choose the tuning parameter m (number of splits) to be $\approx \sqrt{p/3}$, where p is the number of predictors, as is recommended by the creators of the random forest method.

```
# variable importance plot; type = 2 because we are doing classification and are interested in the gini index
varImpPlot(rf.penguin, pch = 20, main = "", type = 2)
```



As is evident from the variable importance plot above (as well as the previously plotted trees from 4c), the most significant variables for prediction of penguin species are bill length and flipper length.

Task 5

a)

FALSE, TRUE, TRUE, TRUE

b)

```
# find optimal cost for support vector classifier
cv.svc = tune(svm, species ~ ., data = train, kernel = "linear", ranges = list(cost = 10^seq(-1, 2, 0.2)))

# create the optimal classifier
svc.best <- cv.svc$best.model

# get predictions
svc.pred <- predict(svc.best, test)

# confusion table
svc.misclass <- table(svc.pred, test$species)

# ----- #

# find optimal cost and gamma parameter for support vector machine with radial kernel
cv.svm = tune(svm, species ~ ., data = train, kernel = "radial", ranges = list(cost = 10^seq(-1, 2, 0.2)))

# create optimal SVM
svm.best <- cv.svm$best.model

# get predictions
svm.pred <- predict(svm.best, test)

# confusion table
svm.misclass <- table(svm.pred, test$species)

# optimal parameter for linear case and corresponding error
summary(cv.svc)[1]

## $best.parameters
##      cost
## 6 1.778279

summary(cv.svc)[2]

## $best.performance
## [1] 0

# optimal parameters for radial case and corresponding error
summary(cv.svm)[1]
```

```
## $best.parameters
##      cost gamma
## 117  100  0.01
```

```
summary(cv.svm)[2]
```

```
## $best.performance
## [1] 0
```

We can observe from the above results that the optimal cost for the linear classifier is 1.778279, while the optimal cost and gamma parameter for the radial SVM are 100 and 0.01, respectively. In both cases, training error is zero.

```
# confusion table for linear classifier
svc.misclass
```

```
##
## svc.pred      Adelie Chinstrap Gentoo
## Adelie        42         0        0
## Chinstrap      0        20        0
## Gentoo         0         0       38
```

```
# confusion table for SVM with radial kernel
svm.misclass
```

```
##
## svm.pred      Adelie Chinstrap Gentoo
## Adelie        42         0        0
## Chinstrap      0        20        0
## Gentoo         0         0       38
```

In both cases, all test penguins are classified correctly, so the misclassification error rates are zero. This makes it difficult to decide on a preferred classifier. Although the radial kernel might be more widely used in practice, we would perhaps prefer the simpler model, because the cross validation and training is less computationally expensive, while still yielding the same results as the radial SVM.

Task 6)

```
set.seed(1)
id <- "1NJ1SuUBeb15P8rMSIwm_n3S8a7K43yP4" # google file ID
happiness <- read.csv(sprintf("https://docs.google.com/uc?id=%s&export=download", id), fileEncoding="UTF-8")

#colnames(happiness)

cols = c('Country.name',
         'Ladder.score', # happiness score
         'Logged.GDP.per.capita',
         'Social.support',
         'Healthy.life.expectancy',
```

```

'Freedom.to.make.life.choices',
'Generosity', # how generous people are
'Perceptions.of.corruption')

# We continue with a subset of 8 columns:
happiness = subset(happiness, select = cols)
rownames(happiness) <- happiness[, c(1)]

# And we creat an X and a Y matrix
happiness.X = happiness[, -c(1, 2)]
happiness.Y = happiness[, c(1, 2)]
happiness.XY = happiness[, -c(1)]

# scale
happiness.X = data.frame(scale(happiness.X))

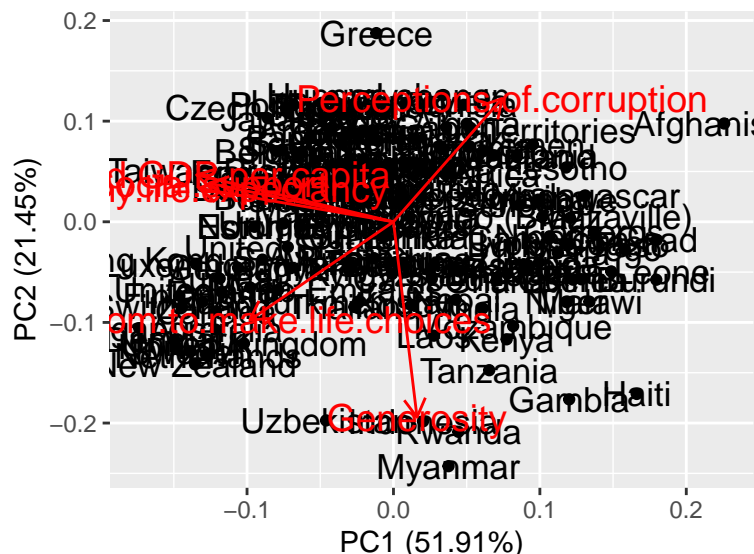
pca_mat = prcomp(happiness.X, center = T, scale = T)

summary(pca_mat)

## Importance of components:
##          PC1      PC2      PC3      PC4      PC5      PC6
## Standard deviation  1.7648  1.1344  0.8384  0.7200  0.50030  0.35657
## Proportion of Variance 0.5191  0.2145  0.1172  0.0864  0.04172  0.02119
## Cumulative Proportion 0.5191  0.7335  0.8507  0.9371  0.97881  1.00000

# Score and loadings plot:
autoplot(pca_mat, data = happiness.X, colour = "Black", loadings = TRUE, loadings.colour = "red",
         loadings.label = TRUE, loadings.label.size = 5, label = T, label.size = 4.5)

```



a)

- i) We are to comment on two characteristics, and the first one we notice is that the predictors “perception of corruption” and “freedom to make life choices” are opposites. This can be interpreted to mean that

in countries with a high score of “perception of corruption” have little “freedom to make choices” and vice versa. This can be seen as they are negatively correlated in both PC1 and PC2. The second characteristic we choose to mention, is that of “generosity” and “healthy life expectancy”. We notice that “generosity” almost only depends on PC2, while “healthy life expectancy” almost only depends on PC1. That would mean that these are nearly uncorrelated. One might then say that “healthy life expectancy” does not affect a countries “generosity” so two countries could be equally generous, even if one has a bad and one has a good “healthy life expectancy” score.

ii) Afghanistan

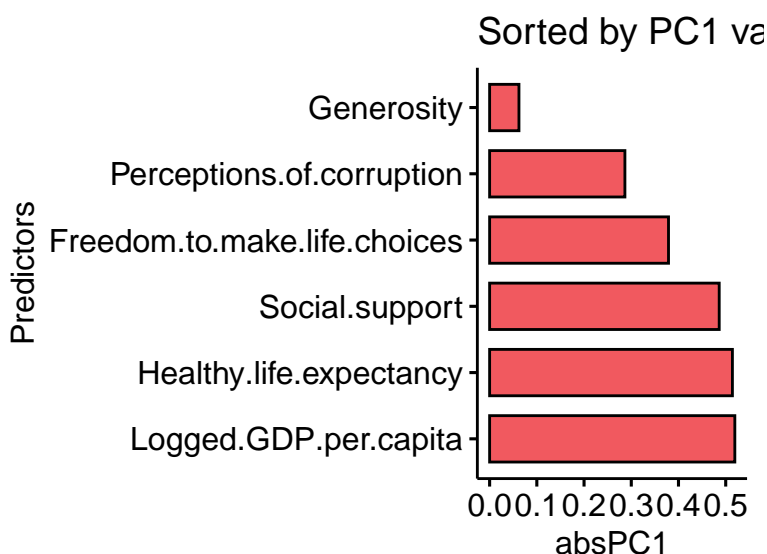
b)

i)

```
set.seed(1)
df<-data.frame(pca_mat$rotation)$PC1

absPC <- data.frame(absPC1 = abs(df),
  Predictors = c(colnames(happiness.XY)[-1]))

ggbarplot(absPC, x = "Predictors", y = "absPC1",
  main="Sorted by PC1 value",
  fill="#f1595f",
  sort.val = "desc",
  sort.by.groups = FALSE,
  x.text.angle = 0,
  orientation="horizontal",
  lab.size = 0.1
)
```



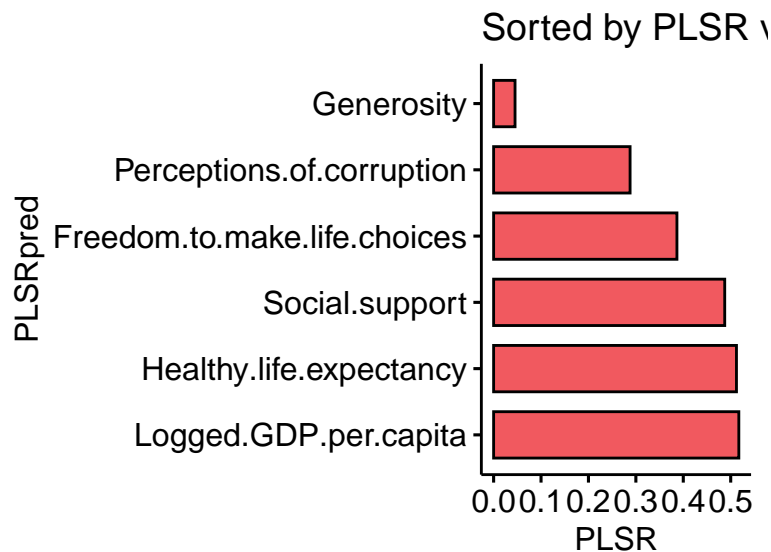
ii)

```
set.seed(1)
plsr_model <- plsr(Ladder.score~., data=happiness.XY, scale=T)
```

iii)

```
set.seed(1)
absPLSR <- data.frame(PLSR =
  abs(plsr_model$loadings[,c('Comp 1')]),
  PLSRpred = c(colnames(happiness.XY)[-1]))

ggbarplot(absPLSR, x = "PLSRpred", y = "PLSR",
  main="Sorted by PLSR value",
  fill="#f1595f",
  sort.val = "desc",
  sort.by.groups = FALSE,
  x.text.angle = 0,
  orientation="horizontal",
  lab.size = 0.1
)
```



As expected, the absolute values of the PCA and the results from the PLSR are very similar.

- iv) The three most important factors based on the PLSR are the predictors corresponding to the largest values in the barplot, i.e. Logged GDP per capita, healthy life expectancy and social support in that order. This is also seen from the PCA barplot, which is very similar to the PSLR barplot.

c)

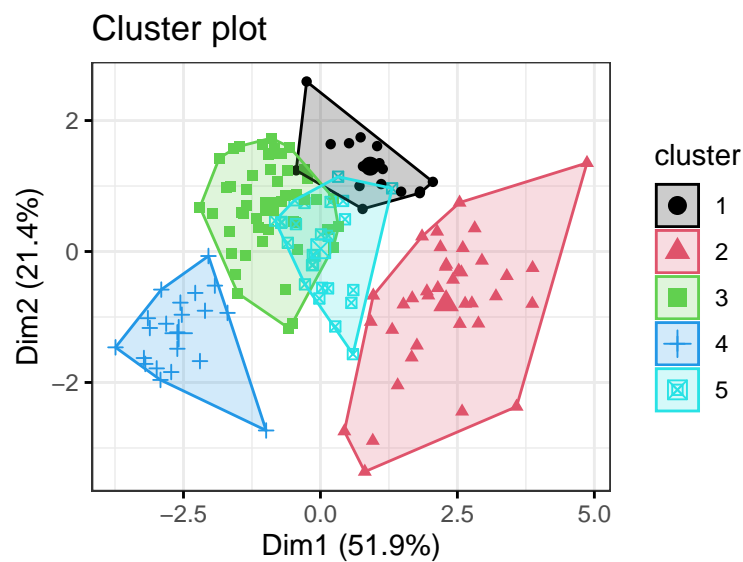
FALSE, FALSE, FALSE, TRUE

d)

i)

```
set.seed(1)
K = 5 # We choose this
km.out = kmeans(happiness.X, K)

fviz_cluster(km.out, data=happiness.X,
              palette = palette(rainbow(5)),
              geom = "point",
              ellipse.type = "convex",
              ggtheme = theme_bw())
```



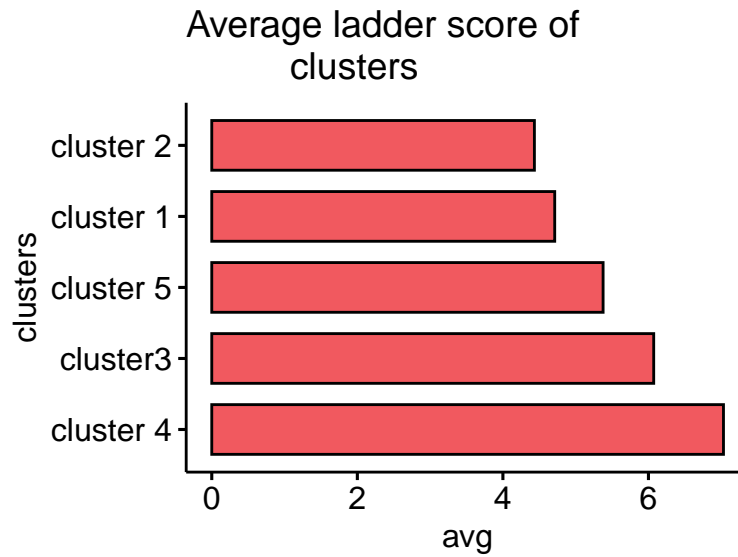
```
cluster_happiness<-function(){
  avgtest=rep(NA, K)
  for (i in 1:K){
    test=km.out$cluster==i
    index=which(test==TRUE)
    avgtest[i]<-mean(happiness.XY[index, ]$Ladder.score)
  }
  return (avgtest)
}

avg_happiness_cluster=cluster_happiness()

avg_df<-data.frame(avg=avg_happiness_cluster,
                   clusters=c("cluster 1", "cluster 2",
                              "cluster3", "cluster 4", "cluster 5"))

ggbarplot(avg_df, x = "clusters", y = "avg",
           main="Average ladder score of
clusters",
           fill="#f1595f",
```

```
sort.val = "desc",  
sort.by.groups = FALSE,  
x.text.angle = 0,  
orientation="horizontal"  
)
```



- ii) We make a plot that is easier to interpret, and we see from the barplot that cluster 4(Finland, Switzerland, Australia) has the highest average ladder/happiness score and cluster 2(Burundi, India, Myanmar) has the lowest. It is also worth to mention that happiness seems to go from left to right in declining order.